

AutoPrompt vs Greedy Coordinate Gradient

Mathis Embit

April 17, 2024

Abstract

We detail two prompt optimization algorithms: AutoPrompt (Shin et al. [2]) and Greedy Coordinate Gradient algorithms (Zou et al. [3]). We show off the differences between the two by unifying the two papers notations, which is especially useful for the gradient computation explanations.

1 Algorithms

1.1 AutoPrompt

Notations:

- x_{inp} : input of the user, which we should not modify
- $x_{trig} = [T] \dots [T]$: triggers tokens we optimize
- $x_{prompt} = \lambda(x_{inp}, x_{trig}) = x_{inp}[T] \dots [T][MASK]$: final prompt sent to the LLM
- w : embedding of the corresponding token $w \in \mathcal{V}$

Algorithm 1 AutoPrompt

```
1: for  $t = 1$  to nb_steps do
2:   for  $x_{trig}^{(j)} \in x_{trig}$  do
3:     for  $w \in \mathcal{V}$  do
4:       compute  $\nabla_{x_{trig}^{(j)}} \log p(y | \{x_{prompt}, \text{ where } x_{trig}^{(j)} = w\})$ 
5:     end for
6:      $\mathcal{V}_{cand} = \text{top-k}_{w \in \mathcal{V}} (w^T \cdot \nabla_{x_{trig}^{(j)}} \log p(y | x_{prompt}))$ 
7:      $x_{trig}^{(j)} = \underset{w_{cand} \in \mathcal{V}_{cand}}{\text{argmax}} p(y | \{x_{prompt}, \text{ where } x_{trig}^{(j)} = w_{cand}\})$ 
8:   end for
9: end for
```

1.2 Greedy Coordinate Gradient

Notations:

- x_{ins}, x'_{ins} : instructions of the aligned LLM (a meta-prompt)
- x_{inp} : input of the user, which we should not modify
- x_{trig} : triggers tokens we optimize
- $x_{prompt} = x_{ins}x_{inp}x_{trig}x'_{ins}$: final prompt sent to the LLM
- \mathcal{I} : indices of the triggers tokens ($x_{trig} = x_{\mathcal{I}}$)
- $\mathcal{L} : x_{1:n} \in \{1, \dots, V\}^n \mapsto -\log p(x_{n+1:n+H}^* | x_{1:n}) \in \mathbb{R}$: the loss function
- $e_{x_i} = [0 \dots 010 \dots 0] \in \{0, 1\}^V$: one-hot encoding of the token $x_i \in \{1, \dots, V\}$
- nb_steps: number of gradient steps

Algorithm 2 Greedy Coordinate Gradient

```

1: for  $t = 1$  to  $\text{nb\_steps}$  do
2:   for  $i \in \mathcal{I}$  do
3:     compute  $-\nabla_{e_{x_i}} \mathcal{L}(x_{1:n})$ 
4:      $\mathcal{X}_i = \underset{w \in \{1, \dots, V\}}{\text{top-k}} -\nabla_{e_{x_i}} \mathcal{L}(x_{1:n})$ 
5:   end for
6:   for  $b = 1, \dots, B$  do
7:      $\tilde{x}_{1:n}^{(b)} := x_{1:n}$ 
8:      $i \sim \mathcal{U}(\mathcal{I})$ 
9:      $\tilde{x}_i^{(b)} \sim \mathcal{U}(\mathcal{X}_i)$ 
10:   end for
11:    $b^* = \underset{b}{\text{argmin}} \mathcal{L}(\tilde{x}_{1:n}^{(b)})$ 
12:    $x_{1:n} = \tilde{x}_{1:n}^{(b^*)}$ 
13: end for

```

2 Common structure

In order to compare the two algorithms let's replace:

- x_{prompt} by $x_{1:n}$ (and remove $[MASK]$ because it's quite confusing)
- $x_{\text{trig}}^{(j)} \in x_{\text{trig}}$ by $i \in \mathcal{I}$
- $\log p(y|x_{1:n})$ by $L(x_i)$ when dealing with the i -th token
- $\nabla_{e_{x_i}} L(w)$ by $w^T \cdot \nabla_{x_i} \log L(x_i)$
- $x_{n+1:n+H}^*$ by y
- $\{1, \dots, V\}$ by \mathcal{V}

Algorithm 3 AutoPrompt

```

1: for  $t = 1$  to  $\text{nb\_steps}$  do
2:   for  $i \in \mathcal{I}$  do
3:     for  $w \in \mathcal{V}$  do
4:       compute  $w^T \cdot \nabla_{x_i} \log L(x_i)$ 
5:     end for
6:      $\mathcal{X}_i = \underset{w \in \mathcal{V}}{\text{top-k}} (w^T \cdot \nabla_{x_i} \log L(x_i))$ 
7:      $x_i = \underset{w_{\text{cand}} \in \mathcal{X}_i}{\text{argmax}} L(w_{\text{cand}})$ 
8:   end for
9: end for

```

Algorithm 4 Greedy Coordinate Gradient

```

1: for  $t = 1$  to  $\text{nb\_steps}$  do
2:   for  $i \in \mathcal{I}$  do
3:     for  $w \in \mathcal{V}$  do
4:       compute  $w^T \cdot \nabla_{x_i} \log L(x_i)$ 
5:     end for
6:      $\mathcal{X}_i = \underset{w \in \mathcal{V}}{\text{top-k}} (w^T \cdot \nabla_{x_i} \log L(x_i))$ 
7:   end for
8:   for  $b = 1, \dots, B$  do
9:      $\tilde{x}_{1:n}^{(b)} := x_{1:n}$ 
10:     $i \sim \mathcal{U}(\mathcal{I})$ 
11:     $\tilde{x}_i^{(b)} \sim \mathcal{U}(\mathcal{X}_i)$ 
12:   end for
13:    $b^* = \underset{b}{\text{argmax}} \log p(y|\tilde{x}_{1:n}^{(b)})$ 
14:    $x_{1:n} = \tilde{x}_{1:n}^{(b^*)}$ 
15: end for

```

3 Gradient computation

Now that we cleared up the notations let's unveil the gradient computation. But first, why do we need to compute a gradient?

If we want an exact solution when dealing with a finite number of possible values we can simply compute the function output for every possible change in the variable to be optimize. In our case we want to

change the i -th token x_i of the prompt and we have a vocabulary of tokens $\mathcal{V} = \{w_1, \dots, w_{|\mathcal{V}|}\}$. We want to find the best token to swap with the current x_i . We could have just computed

$$\begin{bmatrix} p(y|x_{1:n}) & \text{where } x_i = w_1 \\ p(y|x_{1:n}) & \text{where } x_i = w_2 \\ \vdots \\ p(y|x_{1:n}) & \text{where } x_i = w_{|\mathcal{V}|} \end{bmatrix}$$

and picked the token resulting in the highest probability. However this requires $|\mathcal{V}|$ forward pass. Considering the fact that w is a token, i.e. a subset of a word, $|\mathcal{V}|$ is very large (even if we only consider character swap and the alphabet, 26 forward pass is costly).

3.1 AutoPrompt

How to compute $\nabla_{x_i} \log p(y|x_{1:n})$?

Let's write $L(x_1, \dots, x_n) = \nabla_{x_i} \log p(y|x_{1:n})$ the loss we want to maximize, given a token y to generate.

If we write \tilde{x}_i the token x_i after the swap we want to maximize $L(\tilde{x}_i)$ ($= L(x_1, \dots, \tilde{x}_i, \dots, x_n)$) over the possible \tilde{x}_i .

Aproximation: $L(\tilde{x}_i) \approx L(x_i) + \langle \nabla_{x_i} L(x_i), \tilde{x}_i - x_i \rangle = \langle \nabla_{x_i} L(x_i), \tilde{x}_i \rangle + L(x_i) - \langle \nabla_{x_i} L(x_i), x_i \rangle$

So $\operatorname{argmax}_{\tilde{x}_i \in E} L(\tilde{x}_i) \approx \operatorname{argmax}_{\tilde{x}_i \in E} \langle \nabla_{x_i} L(x_i), \tilde{x}_i \rangle$

We derive by the token embeddings hence it's a classic derivative with respect to continuous variable. We can apply the chain rule and the computation requires one forward and one backward pass of the model.

Let's write D the embedding size. $\nabla_{x_i} L(x_{1:n}) \in \mathbb{R}^D$ gives us the exact direction in the embedding space that increases the most the loss when swapping the i -th token.

We cannot directly make a gradient step because it would gives us an embedding that will probably not be the one of an existing token. Instead we should evaluate the gradient in each $w \in \mathcal{V}$: $w^T \cdot \nabla_{x_i} \log p(y|x_{1:n})$. This costs $|\mathcal{V}|$ times the complexity of one projection layer. It informs us on the likelihood increase when replacing the i -th token by w . As we said we can think of this as the swapping making $L(\tilde{x}_i)$ (or the gap $L(\tilde{x}_i) - L(x_i)$) increase the most. We can also think of choosing the token of the vocabulary that is the most colinear with the gradient.

Comment: what about making a gradient ascent and at each step, before evaluating the new current value of the function take as the iterate the closest token in the embedding space? Is it the same? a simple projection? What if we don't pick the closest one at each step and wait for convergence?

3.2 Greedy Coordinate Gradient

How to compute $\nabla_{e_{x_i}} \log p(y|x_{1:n})$?

In fact we will not derive by the token one-hot encoding because it's a discrete variable. The computation will be exactly the same as in AutoPrompt, we just need to exhibit the embedding matrix multiplication.

If we write $\Phi \in \mathbb{R}^{D \times V}$ the embedding matrix (one column = one embedding), we get the embedding x of a one-hot encoded e by performing the multiplication $x = \Phi \cdot e$.

We have:

$$\begin{array}{ccccccc} \{0, 1\}^{|\mathcal{V}|^n} & \longrightarrow & \mathbb{R}^{D^n} & \xrightarrow{\text{LLM}} & \mathbb{R}^D & \longrightarrow & \mathbb{R} \\ (e_1, \dots, e_n) & \longmapsto & (x_1, \dots, x_n) & \longmapsto & x_{n+1} & \longmapsto & L(e_{1:n}) = L(x_{1:n}) \end{array}$$

At index $i \in \mathcal{I}$, when we are optimizing over x_i/e_i , L only depends on x_i/e_i so we write:

$$\begin{array}{ccccccc} \{0, 1\}^{|\mathcal{V}|} & \longrightarrow & \mathbb{R}^D & \xrightarrow{\text{LLM}} & \mathbb{R}^D & \longrightarrow & \mathbb{R} \\ e_i & \longmapsto & x_i & \longmapsto & x_{n+1} & \longmapsto & L(e_i) = L(x_i) \end{array}$$

Let's write

$$h : e_i \in \{0, 1\}^{|\mathcal{V}|} \mapsto \Phi e_i = x_i \in \mathbb{R}^D$$

and

$$g : x_i \in \mathbb{R}^D \mapsto L(x_i) \in \mathbb{R}$$

$$L = g \circ h \text{ so } dL(e) = dg(x) \cdot J_h(e) \text{ and } \nabla_e L(e) = J_h(e)^T \cdot \nabla_x g(x) = \Phi^T \cdot \nabla_x g(x).$$

Finally, by approximating $\tilde{e}_i \in \{0, 1\}^{|\mathcal{V}|}$ with $\tilde{e}_i \in \mathbb{R}^{|\mathcal{V}|}$, we have:

$$\begin{aligned} \operatorname{argmax}_{\tilde{e}_i \in \mathbb{R}^{|\mathcal{V}|}} L(\tilde{e}_i) &\approx \operatorname{argmax}_{\tilde{e}_i \in \mathbb{R}^{|\mathcal{V}|}} \langle \nabla_{e_i} L(e_i), \tilde{e}_i \rangle \\ &= \operatorname{argmax}_{\tilde{e}_i \in \mathbb{R}^{|\mathcal{V}|}} \langle \Phi^T \nabla_{x_i} L(x_i), \tilde{e}_i \rangle \\ &= \operatorname{argmax}_{\tilde{e}_i \in \mathbb{R}^{|\mathcal{V}|}} \tilde{e}_i^T \Phi^T \nabla_{x_i} L(x_i) \\ &= \operatorname{argmax}_{\tilde{e}_i \in \mathbb{R}^{|\mathcal{V}|}} \langle \nabla_{x_i} L(x_i), \Phi \tilde{e}_i \rangle \\ &= \operatorname{argmax}_{\tilde{x}_i \in \mathbb{R}^D} \langle \nabla_{x_i} L(x_i), \tilde{x}_i \rangle \end{aligned}$$

$\nabla_{e_i} L(e_i) \in \mathbb{R}^{|\mathcal{V}|}$ directly gives us how much the likelihood increases when replacing the i -th token by every possible token of the vocabulary \mathcal{V} . However the gradient computation is strictly the same as AutoPrompt.

Comment: exactly the same thing is done in HotFlip (Ebrahimi et al. [1]).

References

- [1] Javid Ebrahimi, Anyi Rao, Daniel Lowd, and Dejing Dou. Hotflip: White-box adversarial examples for text classification, 2018.
- [2] Taylor Shin, Yasaman Razeghi, Robert L. Logan IV au2, Eric Wallace, and Sameer Singh. Auto-prompt: Eliciting knowledge from language models with automatically generated prompts, 2020.
- [3] Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J. Zico Kolter, and Matt Fredrikson. Universal and transferable adversarial attacks on aligned language models, 2023.